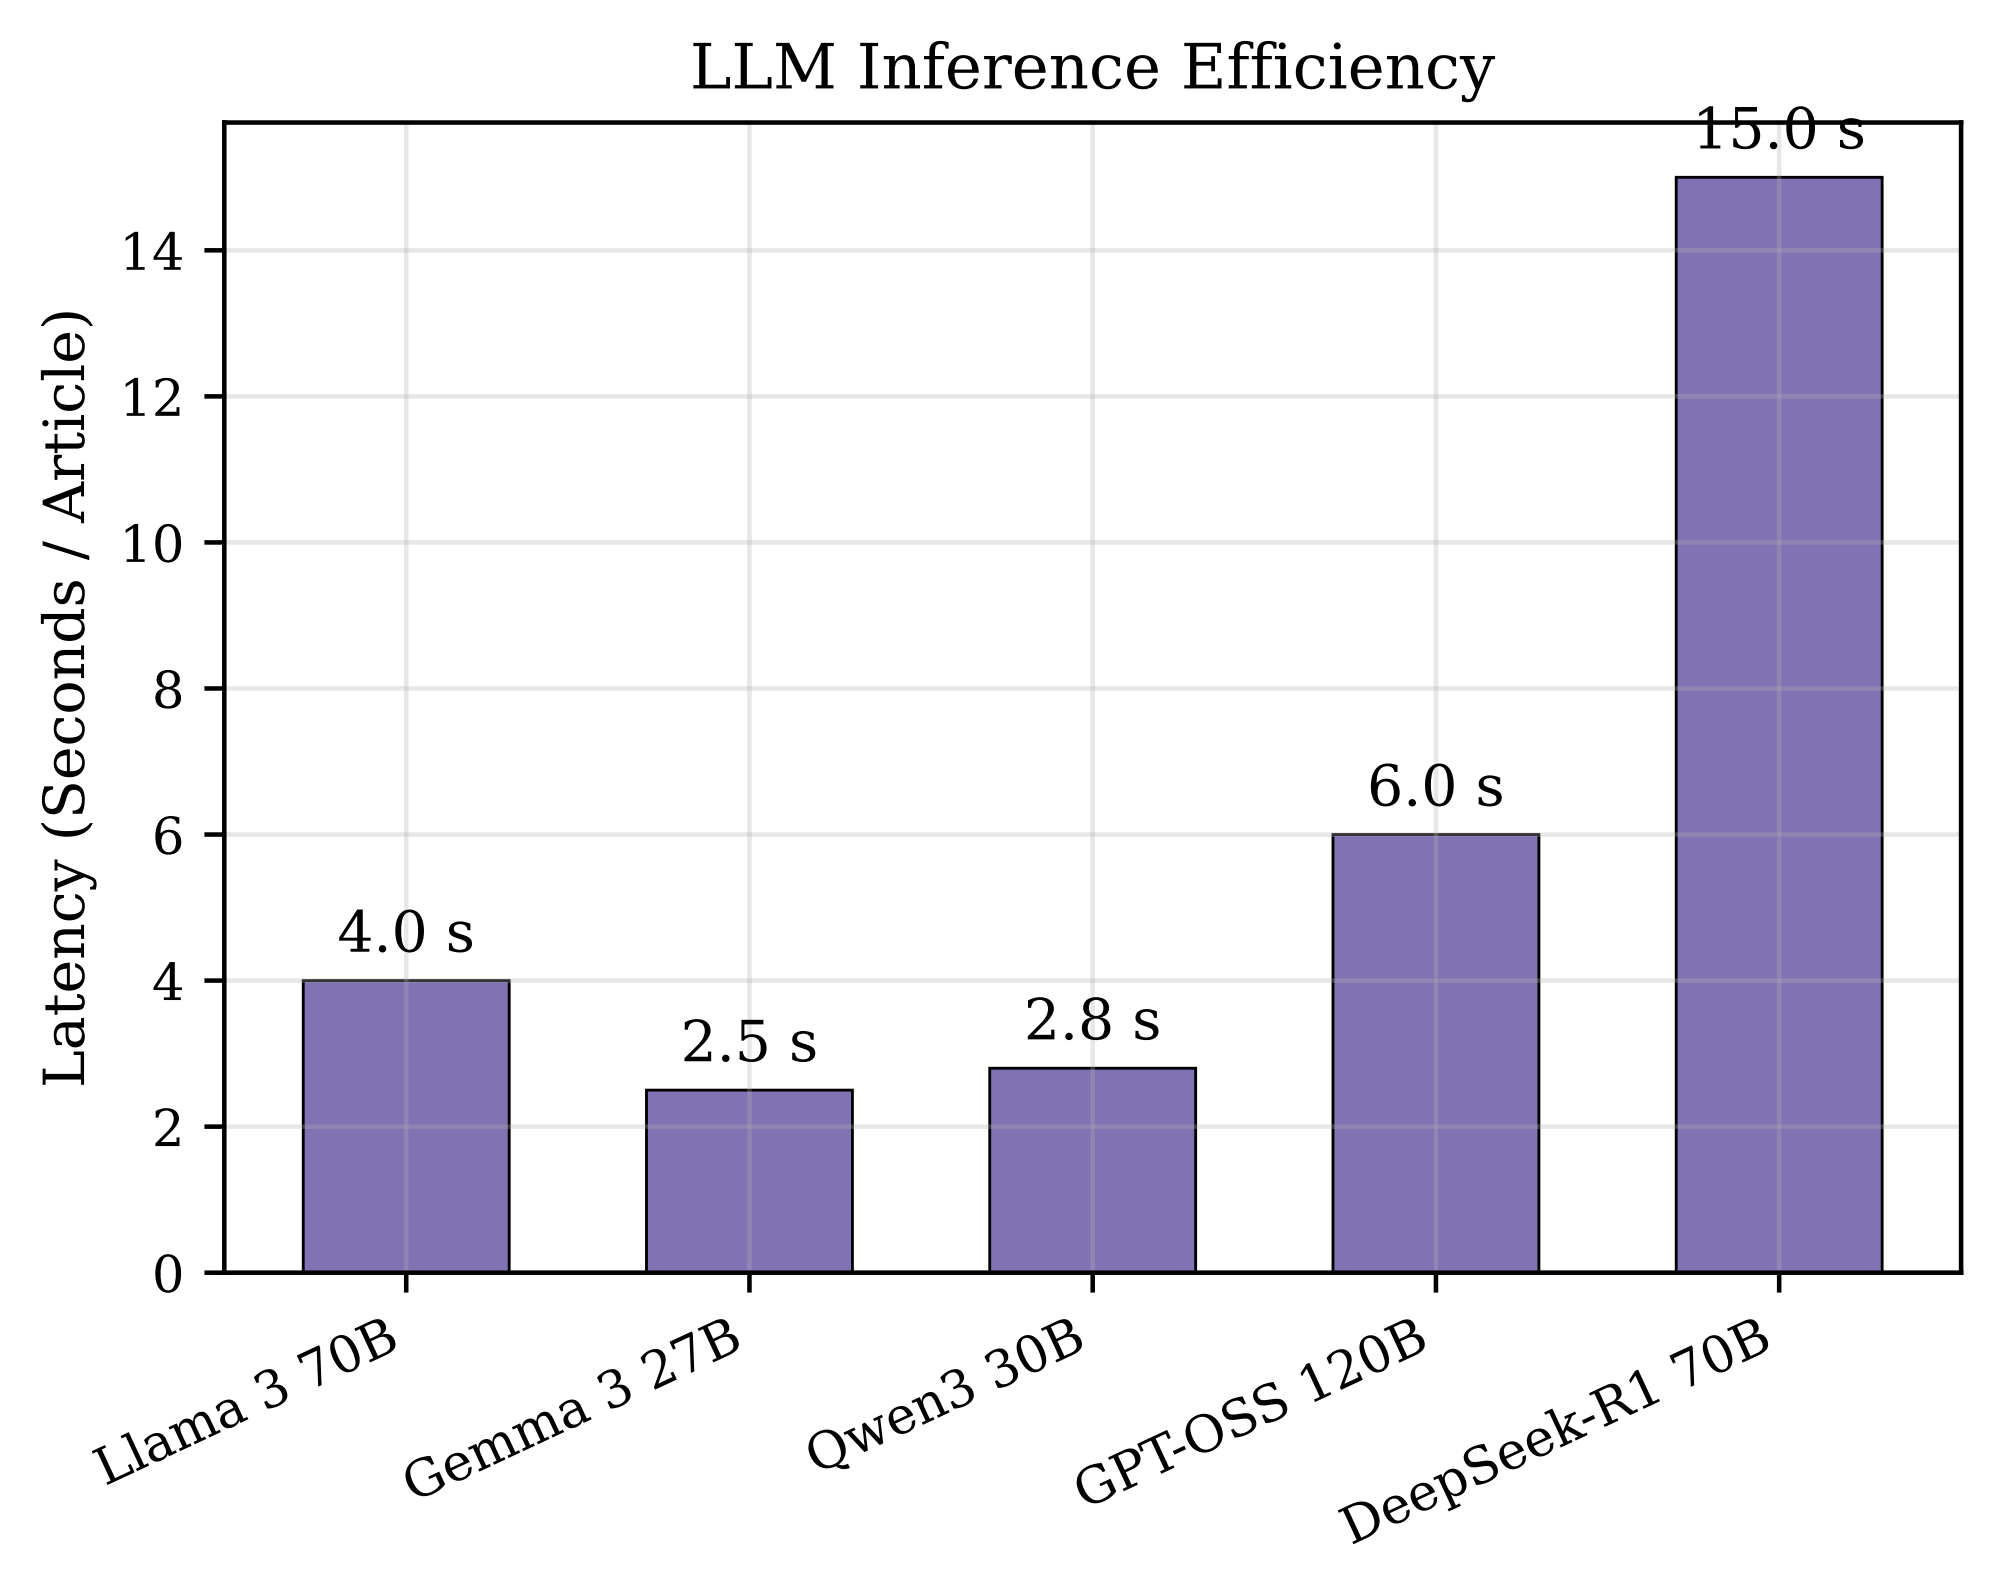
LLM Inference Efficiency
Latency (Seconds / Article)
0
2
4
6
8
10
12
14
4.0 s
2.5 s
2.8 s
6.0 s
15.0 s
Llama 3 70B
Gemma 3 27B
Qwen3 30B
GPT-OSS 120B
DeepSeek-R1 70B